

Strategic Plan for ChefsPlan MVP Launch

Legal Entity and Governance

DAO-Backed Stichting Structure: Establish ChefsPlan as a DAO-governed platform underpinned by a legal **Stichting** (Dutch non-profit foundation). The Stichting will serve as the legal entity to hold assets (e.g. treasury funds, IP) and enter contracts, while the DAO (comprised of token holders) provides community governance. This structure separates **ownership (token-based)** from **legal accountability (foundation board)**. A Dutch *Stichting* has no shareholders and can act in the protocol's interest without profit motive ¹ ². This makes it suitable for a crypto project aiming to be *ownerless* and mission-driven. Consider alternative jurisdictions as well: for instance, Swiss or Liechtenstein foundations are commonly used for crypto projects and offer strong asset protection ³ ⁴. A Swiss foundation (Stiftung) is well-established for non-profit purposes but requires ~CHF 50k endowment and one local board member ⁵ ⁶. Liechtenstein foundations require ~€30k capital and can be non-profit or for-profit, with flexibility and EU market "passporting" access ⁷ ⁸. Weigh these options for regulatory fit; **the Netherlands** offers proximity and familiarity (important for a local MVP launch), whereas **Switzerland** and **Liechtenstein** have explicit legal frameworks for tokenized organizations (e.g. Liechtenstein's Blockchain Act) ⁸. Engaging a legal firm experienced in DAO setups is critical – for example, Swiss firm MME (legal counsel to 1inch DAO ⁹) or Liechtenstein specialists (such as Nägele or Pontinova) – to navigate incorporation and ensure compliance with crypto and labor laws.

Roles: DAO vs. Stichting Board: Define clear governance roles between the token-holder DAO and the foundation's board. The **DAO (token holders)** will ultimately steer major protocol decisions: contract upgrades, fee parameter changes, treasury allocations, and election of certain roles. In the MVP phase, the DAO can exist in a minimal form (e.g. a multisig or simple voting on proposals) while the user base is small. The **Stichting's board** (initially likely the founding team) should have a fiduciary duty to enact the DAO's decisions and handle off-chain responsibilities. Early on, the board will manage critical operations for speed and compliance (signing agreements, paying expenses, handling custody of funds), but it should *cede control progressively* as the DAO community grows ¹⁰ ¹¹. The foundation can also serve as a guardian of last resort – for example, if the DAO vote is blatantly illegal or compromised, the board could have limited veto or delay powers (to be defined in the charter) to protect the project. These powers must be transparent and used only to ensure legal compliance or safety, not to undermine community will. Over time (e.g. after launch phases), move to have the board elected or approved by DAO token holders, solidifying community control while keeping a legal wrapper.

Jurisdiction and Legal Compliance: Given ChefsPlan's focus on chefs and restaurants (initially in the Netherlands), a Dutch Stichting DAO structure is appealing for local compliance. The Netherlands is actively updating rules for foundations and may allow a two-tier board (management and supervisory) ² ¹², which could map to having community representation on a supervisory board. However, Dutch law does not yet formally recognize DAO token voting in corporate governance, so any DAO decisions might be framed as advisory or implemented via board resolutions. To address this, define a **DAO Constitution or Statutes** that commit the board to follow on-chain votes except in narrowly defined situations. In parallel, explore Swiss or Liechtenstein setups if needed for token issuance – e.g. a **Swiss foundation** can provide a tried-and-true model for token projects with regulatory guidance from FINMA on utility vs payment tokens ¹³ ¹⁴, and a **Liechtenstein foundation** offers EU regulatory alignment under MiCA ⁸. Engaging legal partners in the chosen jurisdiction is key. For example, in the Netherlands one might work with firms that have DAO experience to draft the Stichting's articles

(ensuring DAO governance is reflected in foundation bylaws). In Switzerland or Liechtenstein, specialized attorneys (like MME in Zug for Swiss foundations ⁹ or firms familiar with Liechtenstein's TVTG law) can assist. Additionally, ensure compliance with labor and tax laws: since ChefsPlan deals with gig work, the foundation should implement measures to avoid "false self-employment" issues under Dutch law (Wet DBA) ¹⁵ ¹⁶. The **Compliance Agent** (see Infrastructure) and the legal entity together will need to enforce policies (e.g. using model contractor agreements or offering an employment option for frequent gig workers) to keep the platform within legal bounds.

Governance Scope and Processes: For the MVP and beyond, outline how governance proposals and decisions occur. For instance, **protocol upgrades** (smart contract changes) might initially be controlled by a multisig of core team/Foundation (for rapid iteration), with the plan to hand this over to DAO voting or an elected **Tech Committee** by the time of open access. The DAO token (if issued) would be used to vote on upgrade proposals or elect delegates who have contract admin keys. Define a clear upgrade process (e.g. off-chain discussion → on-chain vote → foundation board executes the upgrade transaction as instructed). **Agent permissions:** The platform's autonomous agents (for matching, payouts, etc.) will have admin settings or API keys – initially managed by the team. Governance should specify how these agents get permissions and how they can be replaced or overridden. One approach is to put agent configuration in the hands of the DAO: for example, the DAO could vote to deploy a new version of the Matching agent algorithm or change parameters like fee rates the agent uses. During MVP, these changes can be simple off-chain proposals executed by the team, but with an explicit roadmap to automate them via on-chain governance modules. **Dispute resolution and arbitration** is another area for governance oversight. While day-to-day disputes can be handled by the community staking system (below), the DAO may serve as a final backstop for catastrophic cases. For instance, the DAO could establish an "Arbitration Council" or integrate with an external decentralized court (like Kleros) for appeals. In an extreme scenario (e.g. a dispute that the automated system couldn't resolve), token holders might vote on a resolution or to slash a malicious actor's stake. All such interactions should be codified in the DAO's constitution or governance docs so participants know how protocol updates and disputes are handled.

In summary, **action items** on legal/governance: form the Stichting (or chosen foundation) early, draft statutes aligning with DAO governance principles, allocate initial board roles (likely founding team and perhaps an independent for oversight), and create a lightweight DAO framework (multisig or Snapshot voting) for the MVP phase. As the project matures through launch phases, iterate the governance: expand token holder voting power, onboard community representatives into the foundation or multisigs, and formally transition critical controls from the founding team to the DAO in a stepwise manner ¹⁰ ¹⁷. This phased approach ensures compliance and stability in early days, while laying the groundwork for **full decentralization** when the network grows.

Product and Business Model for MVP

Core MVP Features: The MVP should deliver a focused set of features that solve the key problem – matching chefs ("flexers") with restaurant gigs ("shifts") and ensuring fair payment. We will implement the **core escrow workflow**, a **matching system**, and a **rating/review mechanism** as the backbone of the platform. The escrow workflow means that when a restaurant (planner) hires a chef for a shift, the payment is locked in a smart contract escrow upfront and released upon completion/confirmation of the work ¹⁸. This guarantees chefs get paid on time and builds trust on both sides, much like traditional marketplaces but enforced by code (no manual intervention unless dispute) ¹⁹. The matching system will list open shifts (posted by restaurants) and allow chefs to apply or be algorithmically matched based on criteria (location, skills, rating). Initially, matching can be a simple listing + search interface with perhaps email/app notifications; as a stretch goal, implement an **automatic matching agent** that suggests the best chef for a job based on availability and past

performance. Finally, a basic **rating system** enables both parties to rate each other after a completed shift. These ratings (and possibly written reviews) help build reputation over time. By recording reputation on-chain or in a tamper-proof way, ChefsPlan can ensure credible, portable work history for users ²⁰ – a key advantage over Web2 platforms where reviews can be biased or siloed. Even if the detailed reviews are off-chain, storing a hash on-chain or a cumulative reputation score per user can prevent undetected manipulation and give the DAO a basis to reward top performers in the future.

Monetization Approach: For the MVP, adopt a simple and transparent monetization model that validates the business without creating friction for early users. The recommended approach is a **platform service fee** on each completed gig (escrow release). To stay competitive and mission-aligned, set this fee low relative to traditional staffing agencies or gig platforms – for example, a 5% fee paid by the restaurant per shift payment. Many Web2 freelance platforms charge anywhere from 10% to 20% (often taken from the worker's pay) ²¹ ²², but ChefsPlan can differentiate by charging less and by charging the demand side (restaurants) instead of the chefs. This aligns with emerging Web3 marketplace models: *Braintrust*, for instance, charges clients ~10% while talent keeps 100% of earnings ²³. We can start with a modest fee (e.g. 5-10%) to restaurants, which covers platform costs and provides revenue to the DAO treasury, while chefs receive full pay. This is a **freemium/upside-down model** that makes the platform attractive to chefs (no pay cuts) and still reasonable for restaurants (who are used to agency markups or overtime costs). In addition, consider **staking incentives** in the model: if ChefsPlan has a native token or loyalty points, we could reward users for early participation or good behavior. For example, offer token rewards for the first X shifts completed or for users who stake tokens to support the network (more on staking below). However, token incentives should be used carefully to avoid dependence on speculative rewards; the core revenue should come from delivering value (i.e. successful job matches). A **subscription or premium listing model** could also be explored: perhaps restaurants can pay a monthly subscription for unlimited postings or to get priority placement of their gigs (analogous to “featured job” listings). This could be rolled out after validating basic demand. For MVP simplicity, stick to the per-transaction fee and gather data on willingness to pay.

User Onboarding Strategy: Launching a marketplace requires seeding both sides (chefs and restaurants) in a balanced way. For the MVP, implement a **pilot program** with curated access. Identify a small region or city (e.g. a few major cities in the Netherlands such as Rotterdam/Amsterdam) and onboard an initial cohort of restaurants and freelance chefs. This can be done through personal networks, partnerships, or targeted outreach. For example, the team might partner with a local restaurant association or a culinary school to get a batch of qualified chefs and interested restaurants to trial the platform. **Invite-only onboarding** creates an aura of exclusivity and ensures quality control in early stages. Concretely, we can have a whitelist or referral system: new users sign up by invitation or after a vetting process. This helps maintain trust – early chefs are vetted for reliability, and early restaurants are vetted for reputable practices and prompt payment (though escrow guarantees payment, reputation matters for behavior). During this pilot, closely support users: for instance, assign a team member as an account manager for each pilot restaurant to assist with posting shifts and using the escrow, or provide a quick training to chefs on how to use the app and the importance of confirming shifts. Keeping the community small initially allows for gathering feedback and ensuring each match is successful, which will generate positive testimonials. To attract participants, consider incentives like **zero or reduced fees during the MVP trial**, or even subsidized payments (the platform could cover a small bonus to chefs or a discount to restaurants for first few gigs to encourage usage). These tactics will help jump-start network effects.

Feedback and Iteration Loop: Implement a robust feedback mechanism from day one of the MVP launch. This can include in-app feedback prompts (after a shift completion, ask both parties to rate the experience not just with stars but give qualitative feedback on the platform workflow), regular check-ins with pilot users (e.g. weekly calls or a Telegram/Discord group where users can discuss issues), and

analytics to monitor usage patterns. The team should actively gather improvement suggestions: Is the matching process convenient? Were there any confusions in the escrow or payout? Are the ratings seen as fair? Use this data to drive a rapid iteration cycle. Since the MVP is limited in scope, the development team can push frequent updates – even weekly – to fix bugs or UX pain points uncovered by pilot users ²⁴ ²⁵ . For example, if feedback shows that chefs need a better notification when a shift is approved, prioritize adding SMS or push notifications. If restaurants find the escrow confirmation confusing, improve the UI/UX around payment and release. By the end of the invite-only MVP phase, we should have incorporated the most critical feedback into the product. **Publicly document** the changes and show the community that their input leads to action – this fosters trust and engagement. Finally, set up a channel for ongoing suggestions (like a community forum or governance portal) so that as we open to more users, we can continue the build-measure-learn loop. This iterative mindset not only improves the product but also lays the groundwork for community-driven development, which is essential for a DAO project long term.

Infrastructure Setup

Decentralized Storage (IPFS): ChefsPlan will handle a variety of data – user profiles, job postings, shift agreements, proof of work completion (e.g. a photo of a finished event or a QR code scan), and dispute evidence. For the MVP, set up an **IPFS node or cluster** to store and distribute this content in a tamper-resistant way. Using IPFS (InterPlanetary File System) allows content addressing: each file or piece of data gets a hash, which can be stored on-chain or in our database as an immutable reference. We will likely run our own IPFS **pinning node** on a cloud provider to ensure high availability for MVP ²⁶ ²⁷ . This means deploying an IPFS daemon (or using a service like Infura, Pinata, or web3.storage for convenience) and programming our backend to pin important files. For example, when a shift is completed, the platform could generate a small JSON record of the shift summary (hours worked, agreed payment, any notes) and perhaps a signed acknowledgment from both parties; this JSON and any attached proof (image, PDF) would be stored on IPFS, and the resulting content hash (CID) is what gets referenced in dispute resolution. Early on, it's acceptable that the team's infrastructure is the primary pinner of IPFS content (to guarantee it doesn't vanish). As we progress, we can encourage more nodes to pin data or consider a more decentralized storage incentive (maybe users who run a node get rewarded). IPFS gives us integrity (hashes) and openness, which is superior to keeping all data in a private database – it ensures that if a dispute arises, everyone is referring to the exact same data by hash. For speed and simplicity, smaller metadata (like user reputation scores or shift state hashes) can also be stored on-chain or in an L2 contract, but large blobs (images, detailed logs) will reside in IPFS due to cost.

Layer-2 Smart Contracts (zkSync): To handle transactions (escrow payments, potential staking, etc.) efficiently, we plan to deploy on **zkSync**, a Layer-2 rollup. zkSync's low fees and fast finality make it ideal for an MVP where potentially many micro-transactions (each gig's payment) occur without burdening users with high gas costs ²⁸ ²⁹ . Set up the necessary smart contracts on zkSync testnet first, then zkSync mainnet when ready. Core contracts will include an **Escrow contract** (holds funds for a shift, releases to chef when authorized by both parties or by an agent in case of completion), and perhaps a **Token contract** if a ChefsPlan token is introduced (for governance and staking). The escrow could be as simple as a multi-sig or timelock: funds go in from the restaurant, and a designated "payout agent" or an automated function releases to the chef's address when the shift status is marked complete. zkSync also supports account abstraction, which could later let us sponsor transaction fees for users (so chefs or restaurants don't need to hold ETH for gas, improving UX). For MVP, it's acceptable that users interact with L2 via standard wallets – we'll provide guidance on how to deposit a stablecoin (likely an ETH L2 stablecoin like USDC or DAI) to zkSync. Alternatively, if the user base is not crypto-savvy, we might abstract this by having the foundation accept fiat and fund the escrow on-chain on their behalf (this would be a custodial approach initially). However, one goal is to **abstract crypto complexity** so the

platform feels like using any app, with the blockchain under the hood ³⁰. We can integrate a wallet solution (like Magic.link or WalletConnect) into the app to simplify sign-in and transaction signing for users.

Autonomous Agents (Early Phase): ChefsPlan's architecture calls for autonomous off-chain agents – e.g. a **Matching agent**, **Payout agent**, and **Compliance agent**. In the MVP, these components can run as cloud services managed by the team. Outline each agent's responsibilities and initial implementation:

- **Matching Agent:** This service monitors new shift posts and the pool of available chefs. It could perform tasks like notifying suitable chefs of new gigs or automatically matching a chef if criteria meet (for instance, if a restaurant posts a gig at short notice and a chef with "auto-accept" is available). Initially, this can be a simple cron job or script that sends notifications (email/SMS) to all chefs within X km of the restaurant. More sophisticated matching (algorithmic ranking) can be added as we gather data on what makes a "good match" (skills, rating, past work with that restaurant, etc.). The agent might eventually use AI or optimization algorithms, but MVP can be rule-based. This service runs on a server and uses the API/DB to fetch shifts and users; it doesn't need to be decentralized at first.
- **Payout Agent:** This agent interfaces with the escrow smart contracts on zkSync. It will listen for an event like "ShiftMarkedCompleted" (which could be triggered in the app when both parties confirm completion) and then call the escrow contract to release funds. Alternatively, if a shift completion is time-locked (auto-release after 24 hours if no dispute), the payout agent ensures the transaction occurs on-chain at the right time. For MVP, implement this as a simple Node.js or Python script using web3 libraries, running on a schedule or event subscription, with the Foundation's wallet as the executor. We'll provide it the necessary keys to trigger payouts. This ensures a smooth user experience (chefs get paid promptly) without making each user manually call the contract.
- **Compliance Agent:** Given the regulatory considerations (e.g. verifying that the platform usage doesn't violate labor laws or tax rules), a compliance service is important. At MVP, this might involve checks like: ensuring no single chef works too many hours for one restaurant in a way that looks like full-time employment (which could trigger labor law issues) ¹⁵ ¹⁶; ensuring all users provide required KYC if needed (perhaps verifying identity for big payouts to prevent fraud/money laundering); and monitoring content (profile info, postings) for violations of terms. The compliance agent could be partly manual initially – e.g. an admin dashboard where the team reviews flagged items. However, we can automate some rules: for instance, if a chef has worked >X shifts for the same restaurant in a short period, flag it for review or require that subsequent gigs go through an employment contract (the foundation could even step in to offer a short-term employment contract to be compliant, if that's part of the hybrid model). This agent ensures that as we scale, the platform remains within legal bounds and that **the DAO's operations are socially compliant**, which will be crucial for longevity.

Each of these agents will run on reliable cloud infrastructure (such as AWS, Azure, or a decentralized cloud if feasible). **For an early-stage team, using a standard cloud VM or container service is the fastest route** – e.g. deploy agents as Docker containers that connect to our database and the zkSync network. Document the setup of each (to ease future community operation). We should also instrument these services with logging and alerts; if an agent fails (e.g. payout agent goes down), the team is notified to fix it quickly, as these are critical for user trust.

Progressive Decentralization of Agents: Although agents are centralized in MVP, plan for how to decentralize or distribute them over time. One strategy is to **open source** the agent code and allow community members or partners to run validator or executor nodes that perform the same functions. For example, the matching algorithm could later be run in a way where multiple nodes propose matches and perhaps stake on their proposals' success (creating a competitive marketplace for matching). The payout agent function could eventually be trustless if smart contracts themselves handle more logic (like a fully on-chain escrow with timeouts means no agent needed to trigger). Compliance could move to a community moderation model or integrate with external oracle services

for verification. A concrete step after MVP is to gradually **move key functions on-chain** where possible: e.g., incorporate an on-chain state machine for shifts (open -> filled -> completed -> disputed) that automatically releases payment unless a dispute flag is raised. That reduces reliance on the off-chain payout agent. Similarly, use oracles or **decentralized identity** solutions for compliance (like integrating with a KYC/credential network rather than manual checks). In the interim, the foundation or core team can also decentralize control by using multisigs for agent operations – e.g. the keys that an agent uses to trigger payouts could be held in a 2-of-3 multisig with one key given to a reputable third party or community representative. This way, no single entity can maliciously divert funds, even in the early phase.

Lightweight Staking & Verification (Dispute Infrastructure): Alongside the above, we will implement a preliminary **community verification system** to handle dispute resolution (detailed in the next section). From an infrastructure perspective, this means deploying a **Staking smart contract** on zkSync and integrating it with IPFS and the app. The staking contract will allow certain users (validators) to lock up tokens (or ETH/stablecoin as a stake) which they risk in case of dishonest actions. The system will use **Merkle tree** commitments for data integrity: for each shift, the platform (or the participants) can publish a Merkle root of relevant data (such as clock-in/out times, any evidence files stored on IPFS). The inclusion proofs for pieces of data (like “chef uploaded a photo proof of work”) can be verified by anyone with the Merkle root. The infrastructure needs to support: generating those Merkle roots (likely done off-chain by our backend when a shift is completed and all info is collected), storing the root (perhaps on-chain in an event or in the dispute contract), and a mechanism for validators to retrieve the data from IPFS for checking. We should set up a simple **verification dashboard** where community volunteers can see open disputes or items to verify, download the data via IPFS, and then interact with the staking contract (e.g. by signing a message or calling a function that they attest the data is valid or invalid). This will probably require a web frontend component in the DApp specifically for dispute resolution, but it can be rudimentary in MVP (even a command-line script or a separate web page for testers who are validators). The important part is that the plumbing is in place: IPFS hashes, Merkle proofs, and a contract to accept stakes and slash/reward them based on outcomes. Initially, we might run a **community staking pilot** with a very small group of known community members or even team members pretending to be independent validators, just to test the flows. This can run on testnet in parallel to MVP to gather data on how often disputes happen and how user-friendly the verification process is. Once refined, this system will be a cornerstone of ChefsPlan's trust model, allowing us to handle disputes **without a centralized arbiter**, thereby increasing fairness and scalability ³¹ ³² .

Verification/Dispute Resolution via Community Staking

Even with escrow and ratings in place, disputes can arise (e.g. a restaurant claims a chef left early, or a chef claims the working conditions violated terms). ChefsPlan will implement a **community-driven dispute resolution system** that leverages cryptographic verification and economic incentives to resolve conflicts fairly. The design is inspired by decentralized “optimistic” verification mechanisms and juried DAO courts like Kleros, adapted for our off-chain work context.

Dispute Process Overview: When a shift is completed, if either party **flags a dispute** (for instance, by not confirming completion or by filing a complaint within a certain time window), the dispute resolution process kicks in. The first step is to collect evidence. Both parties can upload relevant evidence to IPFS – e.g. clock-in logs, photos, communication records. The platform's backend then creates a **Merkle tree** of all submitted evidence and relevant data (including perhaps data the platform automatically recorded, like GPS check-in or the original shift agreement). The root hash of this Merkle tree serves as a concise commitment to all evidence. This **Merkle root** could be posted on-chain (to a dedicated Dispute contract) or at least recorded in our database with a timestamp and digital signatures from the parties or platform. Now, to resolve the dispute, **community validators** are called to action.

Community Validator Staking: A pool of trusted community members (eventually open to any DAO token holder meeting certain criteria) can opt in as **dispute validators**. To participate, a validator stakes a certain amount of tokens (or ETH) in the staking contract, signaling they are willing to be chosen as an arbiter. When a dispute arises, a subset of these stakers will be tasked with verifying the evidence. In our model, this isn't a subjective vote on who's right, but rather an **objective verification of facts**: for example, did the hash of an uploaded file match the hash in the Merkle root? Was the "shift completed" event included in the root or omitted? Each validator can independently fetch the evidence from IPFS using the content hashes and compute the Merkle root to confirm it matches the one published ³³. If the evidence clearly supports one side (e.g. a photo with a timestamp proving the chef was present), the validators are essentially checking that and possibly making a judgment call if needed (some disputes might still require interpretation, but we try to minimize that by focusing on factual data). Validators then **cast their judgment** by signing a transaction or voting through the dispute contract: effectively either "Chef's claim is valid" or "Restaurant's claim is valid" (or in simpler terms, yes the work was done satisfactorily, or no it wasn't).

Optimistic Verification and Challenges: To make the system efficient, we use an **optimistic approach**. This means we assume honest behavior by default and only escalate if there's a challenge. For instance, one party (the claimant) could post a proposed resolution (e.g. "release full payment to chef") along with the Merkle root of evidence supporting it. If no one challenges this within a defined period, the resolution goes through automatically. However, if the other party or any community validator believes this is wrong, they can **challenge** by staking some tokens and providing a specific counter-proof (for example, "the chef's photo is missing from the evidence" or "the chef left early, see chat log"). At that point, the system enters a **formal dispute phase** where multiple validators review the evidence. This approach mirrors the mechanism used in optimistic rollups and Truebit: only when a claim is challenged do we execute a detailed verification, saving effort when everyone agrees ³⁴.

Staking Incentives and Penalties: To incentivize honest participation, the system economically rewards correct validators and penalizes dishonest ones. Each validator who stakes to weigh in on a dispute stands to either gain a reward if they align with the truthful outcome or lose part of their stake if they back the wrong side. For example, say a dishonest chef tries to claim they completed a shift without proof. They propose a resolution to pay them and maybe even stake some token as bond. Honest validators would challenge this by providing evidence (or pointing out lack thereof) and staking against the chef's claim. If the evidence doesn't check out (e.g. the Merkle proof shows no sign of completion confirmation), the challenge succeeds – the chef's bond is forfeited, and the challenging validators split that as a reward for keeping the system honest ³⁴. Conversely, if someone falsely challenges a truthful claim (maliciously or mistakenly), that challenger's stake would be forfeited to the honest party. This creates a robust incentive: **only challenge when you are confident, and don't vote in favor of a claim unless you believe it's genuine**. A single honest validator can thwart a false claim, as is typical in optimistic systems where one honest actor is enough to enforce correctness ³⁴. Over time, validators who consistently participate honestly could also earn reputation points or increased voting weight. This is akin to Kleros, where jurors stake a native token (PNK) and are selected for cases, with the chance to earn tokens if they vote with the consensus (which reflects truth) ³⁵ ³⁶. ChefsPlan's token (if available) could serve a similar role: staked for the right to adjudicate and to align incentives with long-term honesty.

Resolution and Enforcement: Once the validators reach a decision (e.g. majority vote or a threshold of stake weight), the dispute contract will enforce the outcome. For instance, if the decision is "chef was in the right, pay them," the escrow payout is released to the chef. If the decision sides with the restaurant (chef didn't fulfill), the escrow might refund the restaurant (or pay partial if that's an option) and perhaps slash the chef's reputation or require further action. The validators who voted with the majority (the correct side) receive a reward – possibly a share of the losing side's stake or a fixed dispute fee.

Those on the minority lose a portion of their stake (which funds the reward) ³⁷ ³⁸ . This mechanism, known as a Schelling game in crypto-economics, encourages voters to assess evidence carefully and aim to be on the truthful side. We will cap the number of validators per dispute initially (to, say, 5 or 7 people) to keep coordination manageable, and we may randomly select them from the pool of stakers to prevent collusion (also similar to how Kleros randomly draws jurors from PNK stakers ³⁹). If a dispute is very high-value or contentious, we can allow an **appeal** process: essentially a larger round of validation with more at stake (again, Kleros does this by increasing juror count on appeal ⁴⁰). In an MVP context, however, most gigs are relatively low-value and we anticipate disputes will be rare and straightforward, so one round with a handful of community verifiers should suffice.

Validator Reputation and Community Trust: To foster a trustworthy pool of validators, the system can maintain a reputation score for each. Validators who have been on the majority (honest) side many times gain reputation, whereas those who consistently lose stakes (indicating poor judgment or malicious intent) could be demoted or even barred. In MVP, since community size is small, we might hand-pick some respected members (or even team members under pseudonyms) to act as initial validators to bootstrap the process. As the user base grows, we'll open this up and perhaps require validators to pass a basic knowledge test of the rules or to have been an active user of the platform for a certain time (to avoid random trolls staking just to grief). We should also clarify that this system is **not meant to replace legal arbitration entirely**; users will agree in Terms of Service that this decentralized resolution is the first recourse. However, part of setting up the legal entity is to ensure that the DAO's dispute process is as legitimate as possible – potentially integrating an off-chain arbitration as a fallback (for example, specifying that disputes not resolved on-chain could be finally settled by an arbitration body under Dutch law if absolutely needed). In practice, we expect the community staking model to handle the vast majority of issues swiftly, without involving courts.

By engaging the community in dispute resolution, ChefsPlan creates a sense of joint ownership of platform integrity. It also reduces the burden on the core team or foundation to mediate every problem. Economically, it aligns incentives so that **truthful behavior is rewarded** and attempting fraud or non-performance is likely to be caught by someone, removing the “free lunch” of cheating the system ⁴¹ ⁴² . This kind of decentralized justice, though experimental, has been tested in other contexts (e.g., Kleros for resolving escrow disputes in services, Aragon Court for governance disputes) and will be a cornerstone in achieving a trust-minimized network.

Deployment Phasing and DAO Onboarding

To mitigate risk and progressively decentralize, we propose launching ChefsPlan in **phases**, each with specific goals and gradually increasing user access and community governance. The phased approach ensures the platform is technically robust and the community is prepared to take on governance as the project matures ¹⁰ ¹¹ .

Phase 0: Devnet and Internal Testing – Before any public launch, deploy the platform on a **development network** (could be a private Ethereum testnet or local zkSync instance) to perform end-to-end tests. Use dummy data to simulate chef and restaurant interactions. This is where the team (and maybe a few advisors) rigorously tests the core flows: posting shifts, escrow payments (with test tokens), completing shifts, triggering disputes, and running the agents. The goal is to catch bugs and ensure the smart contracts and agents behave as expected in a controlled environment. We will also conduct a security review of the contracts at this stage (if resources allow, even a third-party audit for critical escrow logic, given it will hold funds). The devnet phase has no real users or funds at risk, so the team can iterate quickly. Once the application is stable (e.g. MVP features implemented and passing tests ⁴³), we move to a public testnet.

Phase 1: Testnet (Public Beta) – Deploy ChefsPlan on a public testnet (such as zkSync testnet or Goerli) with a limited group of friendly testers. This could include some of the chefs and restaurants from our network who are willing to try the platform with *test tokens*. In this phase, we simulate real usage without financial risk. It allows us to test the crypto-specific components in a real environment: wallet UX, transaction times, IPFS retrieval by different users, and the dispute staking flow using play-money tokens. We'll gather feedback from testers on any confusing blockchain elements. For instance, if chefs struggle with obtaining test ETH to pay gas, we know to improve that flow for mainnet (maybe by auto-funding user wallets with a bit of ETH on sign-up). The testnet phase is also a good time to test the **token distribution mechanics** in a dry run – e.g. we could mint a test token that represents the governance token and practice some voting or distribution (like simulating an airdrop to see how it affects user engagement). Meanwhile, use this phase to start building the DAO community: open a Discord server or forum for testnet participants to discuss and report issues. Educate them on how the planned DAO governance will work. By the end of Phase 1, we should have confidence in the platform's stability and have refined the UX to lower the barrier for non-crypto-savvy users.

Phase 2: Invite-Only MVP Launch (Mainnet) – With lessons learned from testnet, we launch the MVP on zkSync **mainnet** but restrict access to a controlled audience. This invite-only mainnet launch is the true MVP: real chefs, real restaurants, real money in escrow, but gated. We might onboard, say, 10 restaurants and 50 chefs in the first cohort. They will use the platform for actual gigs and transactions. Keeping it invite-only ensures that volume remains manageable and any unforeseen issues can be handled with white-glove support. The DAO governance at this stage can still be *mostly centralized* or council-driven – likely the core team/foundation controls upgrades and handles emergency issues. However, we can start introducing the DAO concept to participants. For example, we could have a simple Snapshot poll (off-chain vote) among early users about a minor policy (like “should we extend the default dispute window from 24h to 48h?”) to get them used to participating in governance, even if the foundation ultimately implements the decision. On the infrastructure side, ensure that the Foundation's multi-sig is set up for key controls (contract ownership, treasury). A multi-sig with 4-of-6 signers including team members and perhaps one trusted external is a good interim step before full DAO control – it adds security/redundancy. Throughout this phase, continue to gather feedback and improve. The goal is to achieve a **solid product-market fit** on a small scale – if the core interactions are smooth and both chefs and restaurants are returning to use it again, we're ready to scale up ⁴⁴ ⁴⁵. If not, we iterate further in invite-only mode until the value prop is clear.

Phase 3: Open Access & Growth – Once the MVP has proven itself with the pilot users (perhaps after a couple of months and a few dozen successful gigs and resolved disputes), remove the invite barrier and open up the platform to the public (initially regionally, then wider). This marks the transition to growth mode. With open access, anyone can sign up as a chef or restaurant and start using ChefsPlan. Accordingly, the **DAO should become more formalized in this phase**. This likely coincides with introducing the governance token (if not already done). We should distribute the token to bootstrap a broad community ownership: a portion could go to the early users (those pilot chefs/restaurants who helped test – essentially a retroactive “airdrop” reward for being first adopters), a portion to the core team/investors (vested over time), and a portion reserved for the foundation treasury or future incentive programs. For example, we might allocate 10% of tokens to early adopters, 20% to team (vested 4 years), 20% to treasury, 50% to be earned by users over time (through usage mining, staking rewards, etc.). If the project will involve fundraising, some of that treasury or a separate allocation could be sold to raise capital (subject to legal compliance with securities laws – which is why a phased approach delaying token launch helps ensure the network is sufficiently decentralized and functional first ¹⁷ ⁴⁶). Technically, in this phase we will also consider deploying on **mainnet Ethereum** or additional L2s if needed for broader reach, but zkSync should suffice initially.

Gradual DAO Governance Onboarding: As the user base grows in Phase 3, governance can be handed off in stages: - **Treasury Control:** Initially, the foundation multi-sig manages funds. Over time, move funds to an on-chain DAO treasury (controlled by token voting). For example, platform fees collected in stablecoins could be periodically transferred to a gnosis-safe wallet where a DAO vote is required to spend them. To start, perhaps require a foundation co-sign (to satisfy legal oversight), but aim to remove that once the community proves responsible budgeting. - **Protocol Upgrade Control:** Implement a timelock + governance module for contracts. For instance, use OpenZeppelin Governor contracts so that token holders can propose and vote on contract upgrades or parameter changes, which execute via a timelock contract. Initially, set high proposal thresholds or allow only whitelisted proposers (like foundation or an elected committee) to avoid governance spam; as more token holders participate, lower these thresholds. - **Agent Administration:** Transition the off-chain agents to community oversight. This could mean publishing the agent source code and allowing tech-savvy community members to run “agent nodes” in exchange for rewards. The DAO could establish an **Agent Operator program** where operators stake some tokens and run the Matching or Compliance agent, and in return get a small cut of platform fees or a periodic token reward. The DAO can then vote to add or remove operators based on performance (similar to how node operators in Livepeer or The Graph are governed). In the interim, the foundation might still run the primary agents, but with increased transparency (publishing logs or metrics to the community). - **Arbitration and Dispute Governance:** With the community staking system in place, the DAO might form a **Dispute Resolution Committee** or simply rely on the stakers and only step in if there's a system failure. One way to involve governance here is for the DAO to vote on parameters like the dispute challenge period duration, the stake amounts, or to select “judges” for an appeals panel. Eventually, the DAO could even integrate a third-party arbitration DAO (for example, making Kleros the final appeal court for very complex disputes) via governance proposal.

At each stage, **educate and involve the community**. For instance, when moving to open access, hold a community call or publish a blog outlining the new governance processes, how token voting works, and encourage users to participate in decision-making. The initial proposals can be about non-critical things to get people comfortable, then ramp up to more important votes (like adjusting fees or rolling out a new feature).

Token Distribution and Launch (if applicable): If a governance token is to be used (likely named something like \$CHEF or \$PLAN), plan its launch carefully. We do not need to launch the token at MVP Day 1 – in fact, it's often wiser to delay token launch until the platform has real usage and a committed community ¹¹. However, since governance and incentives are part of the strategy, we should outline the distribution early to avoid misunderstandings. Possible distribution at mainnet open launch could be: - X% to community treasury (managed by the DAO for future use – grants, liquidity mining, etc.). - Y% to team and early contributors (vested, as noted, to align long-term). - Z% reserved for potential investors or strategic partners (if you plan a token sale or to reward a foundation endowment). - A portion specifically set aside for **user rewards**: e.g., a pool for rewarding validators in disputes (so that when they “earn” something it can be in tokens), a pool for liquidity providers if you need any DeFi integrations, and an airdrop to initial users as mentioned. For example, each pilot chef and restaurant could receive a one-time token grant for being early supporters, which also turns them into DAO voters.

We will likely avoid a wide public token sale until regulatory clarity and the network is truly decentralized (to not run afoul of securities regulations). Instead, if funding is needed, do private SAFT rounds or grants. The token generation event can coincide with a milestone like 100 completed shifts or moving to open access. When launched, list the token on a decentralized exchange on zkSync or Ethereum so it's freely tradeable, providing liquidity for new community members to join the DAO.

Finally, ensure compliance in token distribution: if using a Swiss or Liechtenstein foundation, follow their guidance on utility token classification ¹³ ¹⁴ , and avoid making the token look like a security (i.e., it's for governance and staking, **not** a promise of profit). Progressive decentralization is our shield here: by the time the token is widely distributed, the platform should be demonstrably community-run, which helps from a legal standpoint ¹⁷ ⁴⁶ .

Timeline and Feasibility: An early-stage team can execute this phased approach by focusing on immediate next steps: set up the legal entity and basic DAO framework in parallel with building the MVP. Use Phase 0 and 1 to iron out technical issues (with minimal legal complexity since it's test environment), then formally incorporate before Phase 2 (so that real transactions are under the foundation's umbrella). Each phase acts as a checkpoint to evaluate readiness for the next – if critical issues appear (technical or user adoption), the team can linger in that phase to fix them, rather than rushing ahead. This strategy ensures that when ChefsPlan is fully open and decentralized, it rests on a foundation of proven product-market fit, a growing community, and a compliant legal/governance structure.

By following this strategic plan, ChefsPlan can rapidly deliver an MVP that addresses immediate market needs (connecting chefs with work and restaurants with staff), while simultaneously laying the groundwork for a **decentralized, community-governed network**. Each component – the DAO-backed foundation, the escrow and matching product, the infrastructure and agents, community dispute resolution, and phased deployment – is designed to be *actionable now* for a small team and yet scalable toward the ambitious vision of a self-sustaining protocol. This ensures early success and trust, and positions ChefsPlan to gradually evolve into a fully decentralized marketplace that can operate compliantly and efficiently in the long run.

Sources: The plan leverages best practices and insights from existing projects and research to ensure feasibility: - DAO/Foundation legal structures and jurisdictional considerations ⁵ ⁴⁷
- Freelance marketplace dynamics and the value of blockchain escrow and low fees ²² ²³
- Community-owned network models (e.g. Braintrust's fee structure and non-profit approach) ⁴⁸ ⁴⁹
- Progressive decentralization strategy for handing over control to the community responsibly ¹⁰ ¹¹
- Decentralized dispute resolution mechanisms (optimistic verification, staked jurors) ³⁴ ³⁵ .

¹ Stichting - Wikipedia

<https://en.wikipedia.org/wiki/Stichting>

² Netherlands regulates governance structure for associations and ...

<https://www.pinsentmasons.com/out-law/news/netherlands-regulates-governance-structure-for-associations-and-foundations>

³ ⁵ ⁶ ¹³ ¹⁴ Setting Up a DAO as a Swiss Foundation | Pontinova Law

<https://www.pontinova.law/dao/switzerland-foundation>

⁴ ⁷ ⁸ ⁴⁷ Setting Up a DAO as a Liechtenstein Foundation | Pontinova Law

<https://www.pontinova.law/dao/liechtenstein-foundation>

⁹ 1inch DAO Engages MME as Legal Counsel

<https://www.mme.ch/en/magazine/articles/1inch-dao-engages-mme-as-legal-counsel>

¹⁰ ¹¹ ¹⁷ ⁴⁴ ⁴⁵ ⁴⁶ Progressive Decentralization: A Playbook for Building Crypto Applications |

Andreessen Horowitz

<https://a16z.com/progressive-decentralization-a-playbook-for-building-crypto-applications/>

12 Crack in the Foundation? The Future of Dutch Poison Pills

<https://www.glasslewis.com/article/crack-in-the-foundation-the-future-of-dutch-poison-pills>

15 16 Assessing work relationship between client and contractor (Wet DBA) | Business.gov.nl

<https://business.gov.nl/regulation/employment-relationship-model-agreements-wet-dba/>

18 19 20 21 22 26 27 31 32 7 Must-Have Features for a Blockchain Freelancing Platform

<https://resources.viserlab.com/blockchain-freelancing-platform/>

23 30 41 42 48 49 Braintrust: A Talent Network Governed by its Community | Messari

<https://messari.io/report/braintrust-a-talent-network-governed-by-its-community>

24 25 AGENTS.md

<https://github.com/kickenV/chefsplan5.0/blob/dfdf1722199795e7abc5fb742f384143e18ea369/AGENTS.md>

28 29 34 Ethereum is Changing, and Smart Contract Architecture Should Too

<https://www.archetype.fund/media/ethereum-is-changing-and-smart-contract-architecture-should-too>

33 Sticking to 8192 signatures per slot post-SSF: how and why - Page 2

<https://ethresear.ch/t/sticking-to-8192-signatures-per-slot-post-ssf-how-and-why/17989?page=2>

35 36 38 39 40 PNK Token | Kleros

<https://docs.kleros.io/pnk-token>

37 Blockchain-based Dispute Resolution on the Kleros Platform

<https://journals.library.columbia.edu/index.php/stlr/blog/view/84>

43 instruction_ai.md

https://github.com/kickenV/chefsplan/blob/6c3b92830b492137f5c1e4a9fb5b445cf7f22dda/docs/goals/instruction_ai.md